

Jacob Miller - Homework 3

February 20, 2020

1 Setup

```
[1]: import pandas as pd
import numpy as np
import patsy
import scipy
import statsmodels.api as sm
from astropy.table import Table
from sympy import symbols
```

```
[2]: def title_print(text):
    '''
    Used throughout to print section titles
    '''
    text_len = len(text)
    print()
    print('#' * (text_len + 4))
    print('|', text, '|')
    print('#' * (text_len + 4))
```

2 Problem 3.1

```
[3]: df = pd.read_excel('Data/data-table-B1.xlsx')
df = df.rename(columns = {'y': 'Games_won',
                          'x1': 'Rushing_yards',
                          'x2': 'Passing_yards',
                          'x3': 'Punting_average',
                          'x4': 'Field_goal_percentage',
                          'x5': 'Turnover_differential',
                          'x6': 'Penalty_yards',
                          'x7': 'Percent_rushing',
                          'x8': 'Opponent_rushing_yards',
                          'x9': 'Opponent_passing_yards'})
```

```
[4]: title_print('Problem 3.1a')
y, X = patsy.dmatrices('Games_won ~ Passing_yards + Percent_rushing + \
                        Opponent_rushing_yards', df)

parens = np.matmul(X.T, X)
Xs = np.matmul(np.linalg.inv(parens), X.T)
b_hat = np.round(np.matmul(Xs, y), 4)

print('y_hat = {} + {} * x_2 + {} * x_7 + {} * x_8'.format(b_hat[0],
                                                            b_hat[1],
                                                            b_hat[2],
                                                            b_hat[3]))
```

```
#####
| Problem 3.1a |
#####
y_hat = [-1.8084] + [0.0036] * x_2 + [0.194] * x_7 + [-0.0048] * x_8
```

```
[5]: title_print('Problem 3.1b')
results = sm.OLS(y, X).fit()
results.model.data.design_info = X.design_info

# Note statsmodels prints out ANOVA for each individual regressor
aov_table = sm.stats.anova_lm(results, typ = 1)

print(results.summary())
print('\n--- Analysis of Variance table ---\n{}'.format(aov_table))
print('\nRegression F: {}'.format(round(results.fvalue, 2)))
print('Regression p: {}'.format(round(results.f_pvalue, 4)))
print('\n--> Regression is significant <--')
```

```
#####
| Problem 3.1b |
#####
```

OLS Regression Results			
Dep. Variable:	Games_won	R-squared:	0.786
Model:	OLS	Adj. R-squared:	0.760
Method:	Least Squares	F-statistic:	29.44
Date:	Thu, 20 Feb 2020	Prob (F-statistic):	3.27e-08
Time:	17:27:18	Log-Likelihood:	-52.532
No. Observations:	28	AIC:	113.1
Df Residuals:	24	BIC:	118.4
Df Model:	3		
Covariance Type:	nonrobust		

```

=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept                    -1.8084      7.901      -0.229      0.821     -18.115
14.498
Passing_yards                 0.0036      0.001       5.177      0.000       0.002
0.005
Percent_rushing              0.1940      0.088       2.198      0.038       0.012
0.376
Opponent_rushing_yards      -0.0048      0.001      -3.771      0.001      -0.007
-0.002
=====
Omnibus:                    0.665   Durbin-Watson:                1.492
Prob(Omnibus):              0.717   Jarque-Bera (JB):                0.578
Skew:                      0.321   Prob(JB):                  0.749
Kurtosis:                  2.712   Cond. No.                  7.42e+04
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 7.42e+04. This might indicate that there are strong multicollinearity or other numerical problems.

--- Analysis of Variance table ---

	df	sum_sq	mean_sq	F	PR(>F)
Passing_yards	1.0	76.193400	76.193400	26.172055	3.100132e-05
Percent_rushing	1.0	139.500820	139.500820	47.917840	3.697874e-07
Opponent_rushing_yards	1.0	41.400062	41.400062	14.220716	9.377699e-04
Residual	24.0	69.870004	2.911250	NaN	NaN

Regression F: 29.44

Regression p: 0.0

--> Regression is significant <--

```

[6]: title_print('Problem 3.1c')
t_stat = -scipy.stats.t.ppf(0.025, len(X) - 2)
t_values = abs(np.round(results.tvalues[1:], 3))
p_values = np.round(results.pvalues[1:], 3)
table = Table([[ 'B2', 'B7', 'B8'], t_values, p_values],
              names = ('Coef', 't_0', 'p-value'))

print(table)
print('\nt-statistic = {}'.format(round(t_stat, 3)))

```

```
print('\n--> abs(t_0) > t-statistic, so all are significant <--')
```

```
#####
| Problem 3.1c |
#####
Coef  t_0  p-value
----  ----  -
B2 5.177    0.0
B7 2.198    0.038
B8 3.771    0.001
```

```
t-statistic = 2.056
```

```
--> abs(t_0) > t-statistic, so all are significant <--
```

```
[7]: title_print('Problem 3.1d')
print('R^2 = {}'.format(round(100 * results.rsquared, 2)))
print('Adj-R^2 = {}'.format(round(100 * results.rsquared_adj, 2)))
```

```
#####
| Problem 3.1d |
#####
R^2 = 78.63%
Adj-R^2 = 75.96%
```

```
[8]: title_print('Problem 3.1e')
y2, X2 = patsy.dmatrices('y ~ Passing_yards + Opponent_rushing_yards', df)

parens2 = np.matmul(X2.T, X2)
Xs2 = np.matmul(np.linalg.inv(parens2), X2.T)
b_hat2 = np.round(np.matmul(Xs2, y2), 4)

results2 = sm.OLS(y2, X2).fit()
results2.model.data.design_info = X2.design_info

partial_F = round((results.ess - results2.ess) / results.mse_resid, 2)

print('reduced y_hat = {} + {} * x_2 + {} * x_8'.format(b_hat2[0],
                                                         b_hat2[1],
                                                         b_hat2[2]))

print('partial F = {}'.format(partial_F))
print('\n--> {} < {}, therefore B7 is significant <--'.format(partial_F,
                                                                round(results.fvalue, 2)))
```

```
#####
```

```
| Problem 3.1e |
#####
reduced y_hat = [14.7127] + [0.0031] * x_2 + [-0.0068] * x_8
partial F = 4.83

--> 4.83 < 29.44, therefore B7 is significant <--
```

3 Problem 3.10

```
[9]: df = pd.read_excel('Data/data-table-B11.xlsx')
```

WARNING *** OLE2 inconsistency: SSCS size is 0 but SSAT size is non-zero

```
[10]: title_print('Problem 3.10a')
y, X = patsy.dmatrices('Quality ~ Clarity + Aroma + Body + Flavor + \
                        Oakiness', df)

parens = np.matmul(X.T, X)
Xs = np.matmul(np.linalg.inv(parens), X.T)
b_hat = np.round(np.matmul(Xs, y), 4)

print('y_hat = {} + {} * x_1 + {} * x_2 + {} * x_3 + {} * x_4 + {} * x_5'.\
      format(b_hat[0], b_hat[1], b_hat[2], b_hat[3], b_hat[4], b_hat[5]))
```

```
#####
| Problem 3.10a |
#####
y_hat = [3.9969] + [2.3395] * x_1 + [0.4826] * x_2 + [0.2732] * x_3 + [1.1683] *
x_4 + [-0.684] * x_5
```

```
[11]: title_print('Problem 3.10b')
results = sm.OLS(y, X).fit()
results.model.data.design_info = X.design_info

aov_table = sm.stats.anova_lm(results, typ = 1)

print(results.summary())
print('\n--- Analysis of Variance table ---\n{}'.format(aov_table))
print('\nRegression F: {}'.format(round(results.fvalue, 2)))
print('Regression p: {}'.format(round(results.f_pvalue, 4)))
print('\n--> Regression is significant <--')
```

```
#####
| Problem 3.10b |
```

#####

OLS Regression Results

```

=====
Dep. Variable:          Quality    R-squared:                0.721
Model:                  OLS        Adj. R-squared:            0.677
Method:                 Least Squares    F-statistic:              16.51
Date:                   Thu, 20 Feb 2020    Prob (F-statistic):       4.70e-08
Time:                   17:30:51    Log-Likelihood:           -56.378
No. Observations:       38        AIC:                      124.8
Df Residuals:           32        BIC:                      134.6
Df Model:               5
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	3.9969	2.232	1.791	0.083	-0.549	8.543
Clarity	2.3395	1.735	1.349	0.187	-1.194	5.873
Aroma	0.4826	0.272	1.771	0.086	-0.072	1.038
Body	0.2732	0.333	0.821	0.418	-0.404	0.951
Flavor	1.1683	0.304	3.837	0.001	0.548	1.789
Oakiness	-0.6840	0.271	-2.522	0.017	-1.236	-0.132

```

=====
Omnibus:                1.181    Durbin-Watson:              0.837
Prob(Omnibus):          0.554    Jarque-Bera (JB):          1.020
Skew:                   -0.384    Prob(JB):                  0.601
Kurtosis:               2.770    Cond. No.:                 134.
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

--- Analysis of Variance table ---

	df	sum_sq	mean_sq	F	PR(>F)
Clarity	1.0	0.125210	0.125210	0.092645	7.628120e-01
Aroma	1.0	77.353210	77.353210	57.235072	1.286336e-08
Body	1.0	6.414421	6.414421	4.746149	3.684165e-02
Flavor	1.0	19.049819	19.049819	14.095314	6.945740e-04
Oakiness	1.0	8.597755	8.597755	6.361638	1.683272e-02
Residual	32.0	43.248006	1.351500	NaN	NaN

Regression F: 16.51

Regression p: 0.0

--> Regression is significant <--

```
[12]: title_print('Problem 3.10c')
t_stat = -scipy.stats.t.ppf(0.025, len(X) - 2)
t_values = abs(np.round(results.tvalues[1:], 3))
p_values = np.round(results.pvalues[1:], 3)
table = Table([[ 'B1', 'B2', 'B3', 'B4', 'B5'], t_values, p_values],
              names = ('Coef', 't_0', 'p-value'))

print(table)
print('\nt-statistic = {}'.format(round(t_stat, 3)))
print('\n--> B4, B5: abs(t_0) > t-statistic so these are significant <--')
```

```
#####
| Problem 3.10c |
#####
Coef  t_0  p-value
----  ----  -
B1  1.349   0.187
B2  1.771   0.086
B3  0.821   0.418
B4  3.837   0.001
B5  2.522   0.017
```

```
t-statistic = 2.028
```

```
--> B4, B5: abs(t_0) > t-statistic so these are significant <--
```

```
[13]: title_print('Problem 3.10d')
y2, X2 = patsy.dmatrices('Quality ~ Aroma + Flavor', df)

parens = np.matmul(X2.T, X2)
Xs2 = np.matmul(np.linalg.inv(parens), X2.T)
b_hat2 = np.round(np.matmul(Xs2, y2), 4)
results2 = sm.OLS(y2, X2).fit()
results2.model.data.design_info = X2.design_info

table = Table([[ 'R^2', 'Adj-R^2'],
              [round(100 * results.rsquared, 2),
               round(100 * results.rsquared_adj, 2)],
              [round(100 * results2.rsquared, 2),
               round(100 * results2.rsquared_adj, 2)]]],
              names = (' ', 'Full model', 'Reduced model'))

print(table)
print('\n--> Very similar, so models are similar <--')
```

```
#####
| Problem 3.10d |
```

```
#####
      Full model Reduced model
-----
      R^2      72.06      65.86
Adj-R^2      67.69      63.9
```

--> Very similar, so models are similar <--

```
[14]: title_print('Problem 3.10e')
      ci_1 = np.round(results.conf_int()[4], 3)
      ci_2 = np.round(results2.conf_int()[2], 3)

      print('Full model: {} to {}'.format(ci_1[0], ci_1[1]))
      print('Reduced model: {} to {}'.format(ci_2[0], ci_2[1]))
      print('\n--> Very similar again, so similar models <--')
```

```
#####
| Problem 3.10e |
#####
Full model: 0.548 to 1.789
Reduced model: 0.58 to 1.76
```

--> Very similar again, so similar models <--

4 Problem 3.25 (Note: This is problem 3.21 in 4th edition of textbook)

```
[15]: title_print('Problem 3.25a')
      b, b0, b1, b2, b3, b4 = symbols('b b0 b1 b2 b3 b4')
      y, x1, x2, x3, x4, eps = symbols('y x1 x2 x3 x4 eps')
      gamma_0, gamma_1, z = symbols('gamma_0 gamma_1 z')

      beta = np.array([[b0], [b1], [b2], [b3], [b4]])
      X = np.array([1, x1, x2, x3, x4])
      y = np.matmul(X, beta) + eps

      # H0: b1 = b2 = b3 = b4
      beta2 = np.array([[b0], [b1], [b1], [b1], [b1]])
      y2 = np.matmul(X, beta2) + eps

      T = np.array([[0, 1, -1, 0, 0],
                    [0, 0, 1, -1, 0],
                    [0, 0, 0, 1, -1]])
      c = np.array([[0], [b], [b], [b], [b]])
```



```

print('y = {}'.format(y))
print('H0: b1 = b2 = b3 = b4 = b')
print('\nTherefore: b1 - b2 = 0, b2 - b3 = 0, b3 - b4 = 0')
print('\nT = \n{}\n\nbeta = \n{}\n\nc = \n{}'.format(T, beta, c))
print('\ny = {}'.format(y2))
print('\nWhere:\ngamma_0 = b0\ngamma_1 = b\nz = x1 + x2 + x3 + x4')
print('\n--> Reduced model: y = gamma_0 + gamma_1 * z + eps <--')

```

#####

| Problem 3.25a |

#####

y = [b0 + b1*x1 + b2*x2 + b3*x3 + b4*x4 + eps]

H0: b1 = b2 = b3 = b4 = b

Therefore: b1 - b2 = 0, b2 - b3 = 0, b3 - b4 = 0

T =

```

[[ 0  1 -1  0  0]
 [ 0  0  1 -1  0]
 [ 0  0  0  1 -1]]

```

beta =

```

[[b0]
 [b1]
 [b2]
 [b3]
 [b4]]

```

c =

```

[[0]
 [b]
 [b]
 [b]
 [b]]

```

y = [b0 + b1*x1 + b1*x2 + b1*x3 + b1*x4 + eps]

Where:

gamma_0 = b0

gamma_1 = b

z = x1 + x2 + x3 + x4

--> Reduced model: y = gamma_0 + gamma_1 * z + eps <--

```
[16]: title_print('Problem 3.25b')

beta = np.array([[b0], [b1], [b2], [b3], [b4]])
X = np.array([1, x1, x2, x3, x4])
y = np.matmul(X, beta) + eps

# H0: b1 = b2, b3 = b4
beta2 = np.array([[b0], [b1], [b1], [b3], [b3]])
y2 = np.matmul(X, beta2) + eps

T = np.array([[0, 1, -1, 0, 0],
              [0, 0, 0, 1, -1]])
c = np.array([[0], [0]])

print('y = {}'.format(y))
print('H0: b1 = b2, b3 = b4')
print('\nTherefore: b1 - b2 = 0, b3 - b4 = 0')
print('\nT = \n{}\n\nbeta = \n{}\n\nc = \n{}'.format(T, beta, c))
print('\ny = {}'.format(y2))
print('\nWhere: \ngamma_0 = b0\ngamma_1 = b1\ngamma_3 = b3')
print('z1 = x1 + x2\nz3 = x3 + x4')
print('\n--> Reduced model: y = gamma_0 + gamma_1 * z1 + gamma_3 * z3 <--')
```

```
#####
| Problem 3.25b |
#####
y = [b0 + b1*x1 + b2*x2 + b3*x3 + b4*x4 + eps]
H0: b1 = b2, b3 = b4

Therefore: b1 - b2 = 0, b3 - b4 = 0

T =
[[ 0  1 -1  0  0]
 [ 0  0  0  1 -1]]

beta =
[[b0]
 [b1]
 [b2]
 [b3]
 [b4]]

c =
[[0]
 [0]]
```

```
y = [b0 + b1*x1 + b1*x2 + b3*x3 + b3*x4 + eps]
```

Where:

```
gamma_0 = b0
```

```
gamma_1 = b1
```

```
gamma_3 = b3
```

```
z1 = x1 + x2
```

```
z3 = x3 + x4
```

```
--> Reduced model: y = gamma_0 + gamma_1 * z1 + gamma_3 * z3 <--
```