

# Jacob Miller - Homework 4

March 2, 2020

```
[1]: import pandas as pd
import numpy as np
import patsy
import statsmodels.api as sm
import scipy
import matplotlib.pyplot as plt
from astropy.table import Table
```

## 0.1 Problem 1

```
[2]: # Problem setup
SS_reg = 5550.8166
SS_tot = 5784.5426

table = Table([[ 'Regression', 'Residual', 'Total'], # Source of Variation
               [SS_reg, '[ ]', SS_tot], # Sum of Squares
               ['[ ]', '[ ]', '[ ]'], # Degrees of Freedom
               ['[ ]', '[ ]', ''], # Mean Square
               ['[ ]', ' ', ''], # F0
               ['[ ]', ' ', '']], # P-value
             names = ('Source of Variation', 'Sum of Squares',
                      'Degrees of Freedom', 'Mean Square', 'F0', 'P-value'))

# Sum of Squares
SS_res = round(SS_tot - SS_reg, 4)

# Degrees of Freedom
DoF_reg = 2
DoF_tot = 25 - 1
DoF_res = round(DoF_tot - DoF_reg, 4)

# Mean Squares
MS_reg = round(SS_reg / DoF_reg, 4)
MS_res = round(SS_res / DoF_res, 4)

# F0
```

```

F0 = round(MS_reg / MS_res, 4)

# P-value
P = 1 - scipy.stats.f.cdf(F0, DoF_reg, DoF_res)

# Final table
table = Table(['Regression', 'Residual', 'Total'], # Source of Variation
              [SS_reg, SS_res, SS_tot], # Sum of Squares
              [DoF_reg, DoF_res, DoF_tot], # Degrees of Freedom
              [MS_reg, MS_res, ''], # Mean Square
              [F0, '', ''], # F0
              [P, '', ''], # P-value
              names = ('Source of Variation', 'Sum of Squares',
                       'Degrees of Freedom', 'Mean Square', 'F0', 'P-value'))

```

```
[3]: print(table.to_pandas().to_string())
```

	Source of Variation	Sum of Squares	Degrees of Freedom	Mean Square	F0
0	Regression	5550.8166	2	2775.4083	261.2419
1	Residual	233.7260	22	10.6239	
2	Total	5784.5426	24		

```
[4]: print('--> Small P-value [{:e}] | Reject null hypothesis <--'.format(P))
```


```
--> Small P-value [4.440892e-16] | Reject null hypothesis <--
```

## 1 Problem 2

```
[5]: df = pd.DataFrame(data = {'Pressure_(x)': [30, 31, 32, 33, 34, 35, 36],
                              'Mileage_(y1)': [29.5, 32.1, 36.3, 38.2,
                                                37.7, 33.6, 26.8],
                              'Mileage_(y2)': [30.2, 34.5, 35.0, 37.6,
                                                36.1, 34.2, 27.4]}). \
      set_index('Pressure_(x)')
```

### 1.0.1 Problem 2.a

```
[6]: df_com = df[df.columns[0]].append(df[df.columns[1]])
plt.scatter(x = df_com.index.values, y = df_com)
plt.xlabel('Pressure (psi)')
plt.ylabel('Milage (thousands)')
plt.show()
print('--> Optimal tire pressure appears to be between 32 - 34 psi <--')
```



--> Optimal tire pressure appears to be between 32 - 34 psi <--

### 1.0.2 Problem 2.b

```
[7]: y, X = patsy.dmatrices('df_com.values ~ df_com.index', df_com)

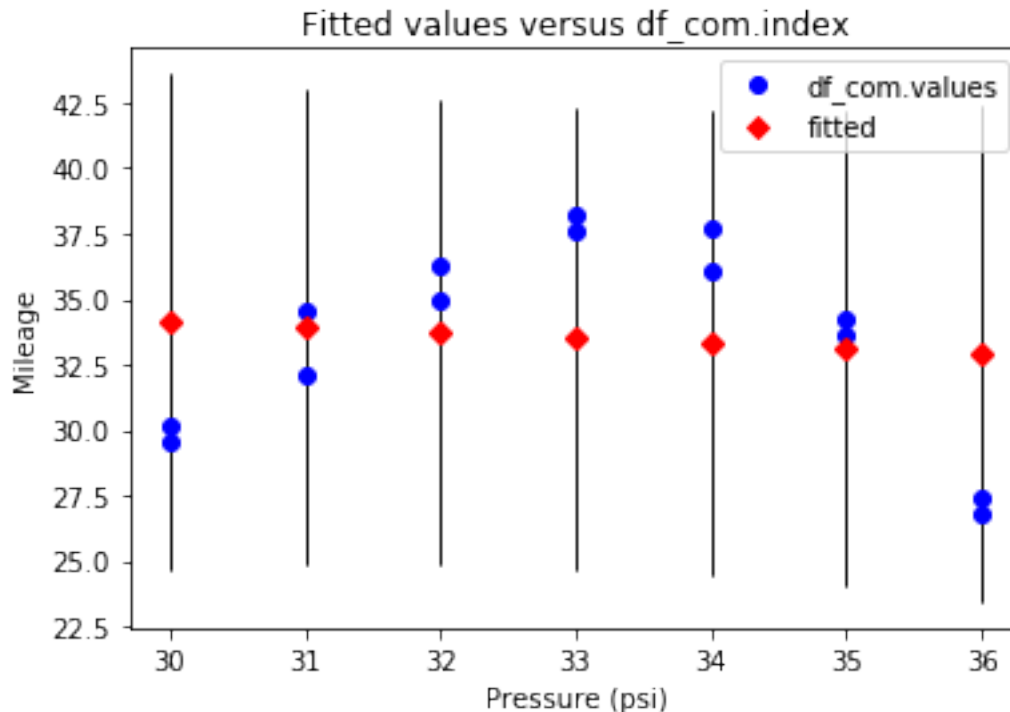
model = sm.OLS(y, X)
results = model.fit()
results.model.data.design_info = X.design_info

print('--> y = {} + {} * x <--'.format(*np.round((results.params[0],
                                                    results.params[1]),
                                                    2)))

fig, ax = plt.subplots()
fig = sm.graphics.plot_fit(results, 1, ax = ax)
ax.set_xlabel('Pressure (psi)')
```

```
ax.set_ylabel('Mileage')
plt.show()
```

-->  $y = 40.35 + -0.21 * x$  <--



### 1.0.3 Problem 2.c

```
[8]: # Setup - because each X has exactly 2 y, can just append to dataframe
df['Mean'] = df.apply(lambda x: np.mean(x), axis = 1)
df['PE Sums'] = df.apply(lambda x: ((x['Mileage_(y1)'] - x['Mean']) ** 2 +
                                     (x['Mileage_(y2)'] - x['Mean']) ** 2),
                        axis = 1)

# Get beta-hat, start to calculate SS_reg
b_hat = np.matmul(np.matmul(np.linalg.inv(np.matmul(X.T, X)), X.T), y)
first_term = np.matmul(np.matmul(b_hat.T, X.T), y)
second_term = sum(y)**2 / len(y)

# Sum of Squares
SS_tot = np.round(float(sum(y**2) - sum(y)**2 / len(y)), 3)
SS_reg = np.round(float(first_term - second_term), 3)
SS_res = np.round(float(SS_tot - SS_reg), 3)
```

```

SS_pe = np.round(float(sum(df['PE Sums'])), 3)
SS_lof = np.round(float(SS_res - SS_pe), 3)

# Degrees of Freedom
DoF_reg = 1
DoF_tot = len(df_com) - 1
DoF_res = DoF_tot - DoF_reg
DoF_pe = len(df) # because each X has exactly 2 y, this is just num of X's
DoF_lof = DoF_res - DoF_pe

# Mean Squares
MS_reg = np.round(float(SS_reg / DoF_reg), 3)
MS_res = np.round(float(SS_res / DoF_res), 3)
MS_pe = np.round(float(SS_pe / DoF_pe), 3)
MS_lof = np.round(float(SS_lof / DoF_lof), 3)

# F0
F0 = round(MS_lof / MS_pe, 3)

# P-value
P = round(1 - scipy.stats.f.cdf(F0, DoF_reg, DoF_res), 10)

table = Table(['Regression', 'Error', 'Lack of Fit', 'Pure Error', 'Total'],
              [SS_reg, SS_res, SS_lof, SS_pe, SS_tot], # Sum of Squares
              [DoF_reg, DoF_res, DoF_lof, DoF_pe, DoF_tot], # Degrees of Freedom
              [MS_reg, MS_res, MS_lof, MS_pe, ''], # Mean Square
              ['', '', F0, '', ''], # F0
              ['', '', P, '', '']], # P-value
              names = ('Source of Variation', 'Sum of Squares',
                      'Degrees of Freedom', 'Mean Square', 'F0', 'P-value'))

```

```
[9]: print(table.to_pandas().to_string())
```

	Source of Variation	Sum of Squares	Degrees of Freedom	Mean Square	F0
0	Regression	2.403	1	2.403	
1	Error	183.434	12	15.286	
2	Lack of Fit	177.644	5	35.529	42.961
2.71169e-05					
3	Pure Error	5.790	7	0.827	
4	Total	185.837	13		

#### 1.0.4 Problem 2.d

```
[10]: print('--> R = {} | Low R-squared suggests model does not fit data well <--'.\
      format(round(results.rsquared, 5)))
```

```
--> R = 0.01293 | Low R-squared suggests model does not fit data well <--
```

#### 1.0.5 Problem 2.e

```
[11]: print('--> P = {} | Reject hypothesis that model describes data <--'.format(P))
```

```
--> P = 2.71169e-05 | Reject hypothesis that model describes data <--
```

#### 1.0.6 Problem 2.f

```
[12]: print('--> Assumed first order linear regression, bad assumption <--')
```

```
--> Assumed first order linear regression, bad assumption <--
```

## 2 Problem 4.3

```
[13]: df = pd.read_excel('Data/data-table-B2.xlsx')
      y, X = patsy.dmatrices('y ~ x4', df)
      model = sm.OLS(y, X)
      results = model.fit()
      results.model.data.design_info = X.design_info
```

#### 2.0.1 Problem 4.3.a


```
[14]: # Get residuals and probability for plot
      residuals = results.resid
      Prob = [(i - 1/2) / len(y) for i in range(len(y))]

      # Calculate OLS using residuals to plot straight line. Get y values from model
      resid_results = sm.OLS(Prob, sm.add_constant(sorted(residuals))).fit()
      X_range = np.linspace(min(residuals), max(residuals), len(residuals))

      # Normal Probability Plot + straight line
      fig, ax = plt.subplots()
      ax.scatter(sorted(residuals), Prob)
      ax.plot(X_range, resid_results.params[0] + resid_results.params[1] * X_range)
      ax.set_xlabel('Residual')
      ax.set_ylabel('Probability')
```

```
plt.title('Normal Probability Plot')
plt.show()

print('--> Minimal fluctuation suggests normality is ok <--')
```

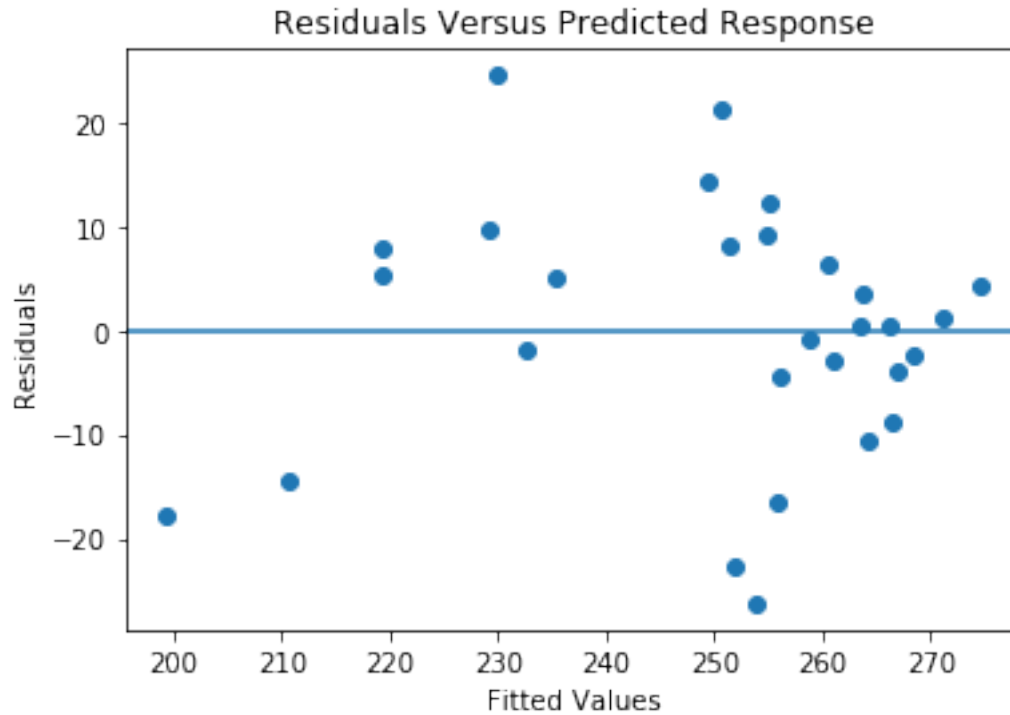


```
--> Minimal fluctuation suggests normality is ok <--
```

## 2.0.2 Problem 4.3.b

```
[15]: fig, ax = plt.subplots()
ax.scatter(results.fittedvalues, residuals)
ax.axhline(0)
ax.set_xlabel('Fitted Values')
ax.set_ylabel('Residuals')
plt.title('Residuals Versus Predicted Response')
plt.show()

print('--> Appears to be non-constant variance (funnel or double-bow) <--')
```



--> Appears to be non-constant variance (funnel or double-bow) <--

### 3 Problem 4.5

```
[16]: df = pd.read_excel('Data/data-table-B4.xlsx')
y, X = patsy.dmatrices('y ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 + x8 + x9', df)
model = sm.OLS(y, X)
results = model.fit()
results.model.data.design_info = X.design_info
```

#### 3.0.1 Problem 4.5.a

```
[17]: # Get residuals and probability for plot
residuals = results.resid
Prob = [(i - 1/2) / len(y) for i in range(len(y))]

# Calculate OLS using residuals to plot straight line. Get y values from model
resid_results = sm.OLS(Prob, sm.add_constant(sorted(residuals))).fit()
X_range = np.linspace(min(residuals), max(residuals), len(residuals))

# Normal Probability Plot + straight line
```




```

fig, ax = plt.subplots()
ax.scatter(sorted(residuals), Prob)
ax.plot(X_range, resid_results.params[0] + resid_results.params[1] * X_range)
ax.set_xlabel('Residual')
ax.set_ylabel('Probability')
plt.title('Normal Probability Plot')
plt.show()

print('--> Minimal fluctuation suggests normality is ok <--')

```



--> Minimal fluctuation suggests normality is ok <--

### 3.0.2 Problem 4.5.b

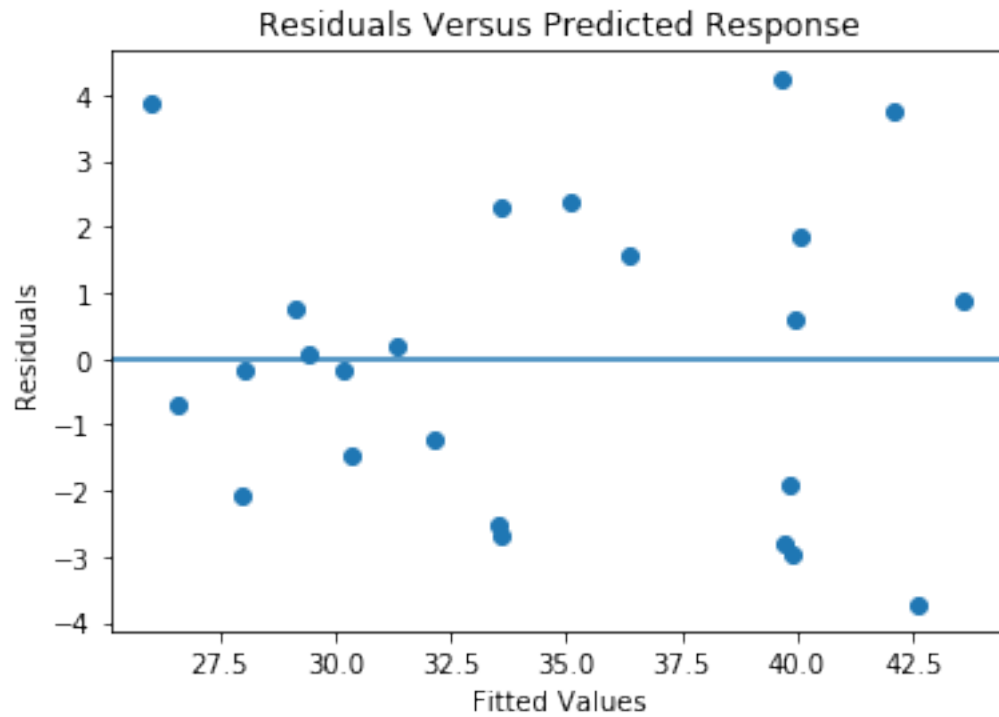
```

[18]: fig, ax = plt.subplots()
ax.scatter(results.fittedvalues, residuals)
ax.axhline(0)
ax.set_xlabel('Fitted Values')
ax.set_ylabel('Residuals')
plt.title('Residuals Versus Predicted Response')
plt.show()

print('--> If outlier in top left corner is removed, plot has incline <--')

```

```
print('--> from lower left to top right. Otherwise no significant pattern <--')
```

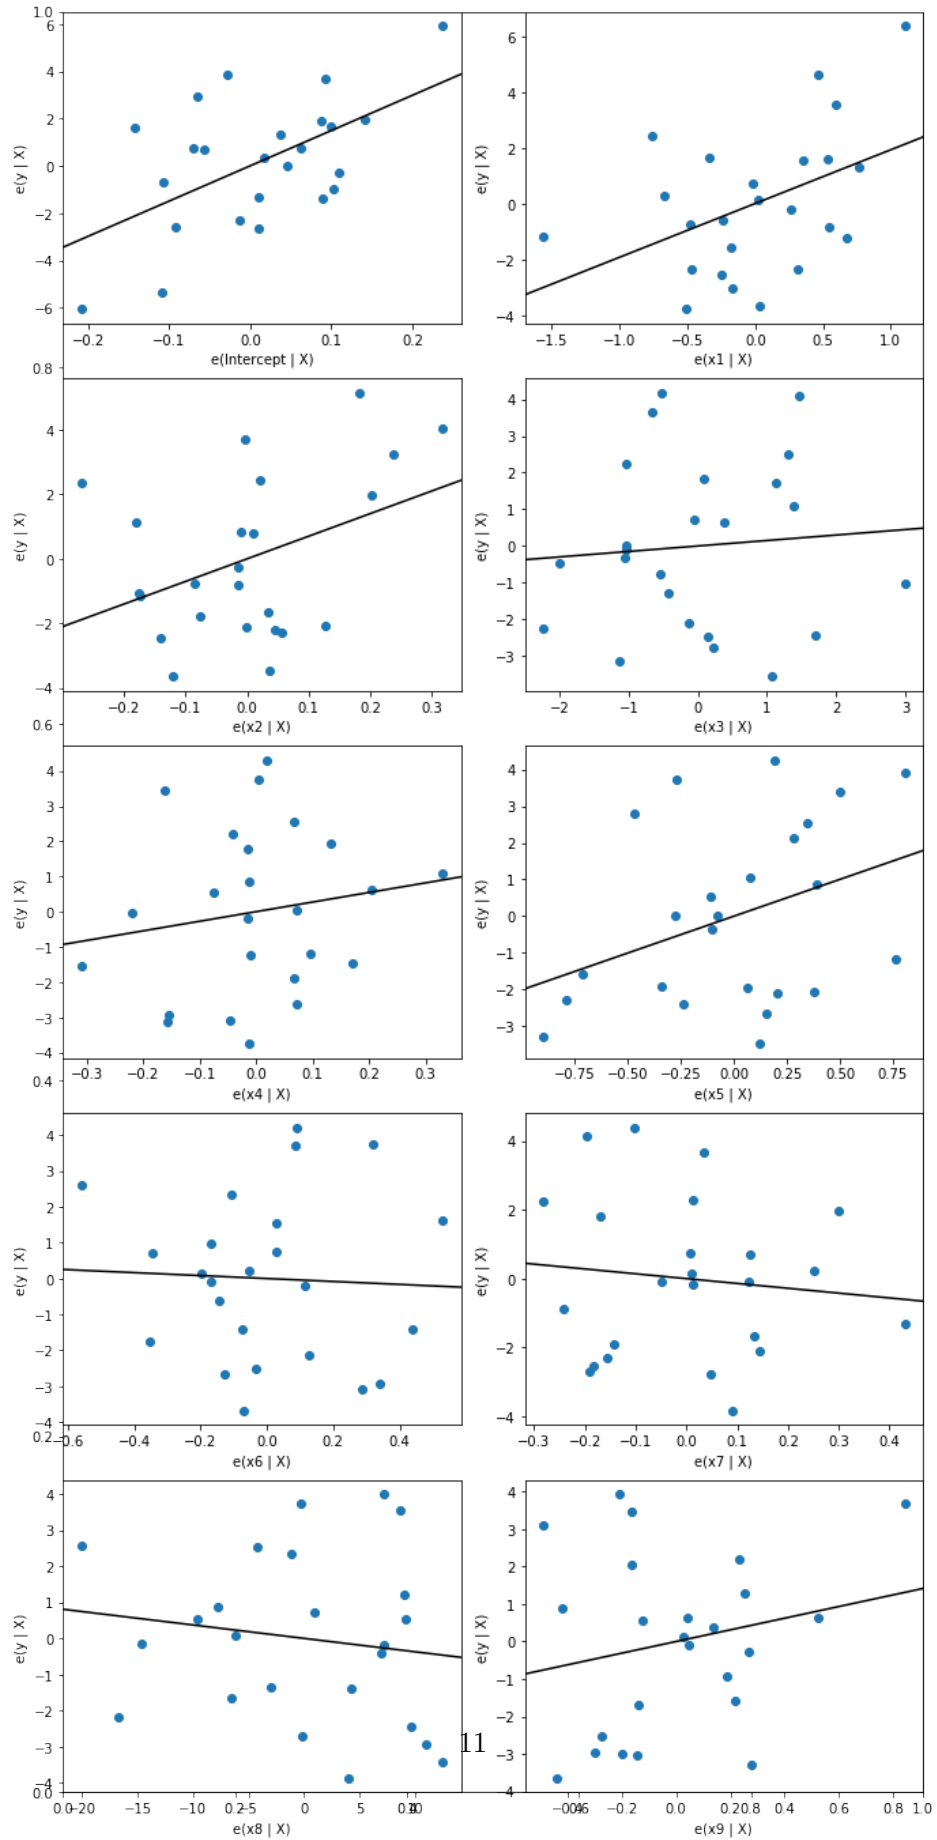


```
--> If outlier in top left corner is removed, plot has incline <--  
--> from lower left to top right. Otherwise no significant pattern <--
```

### 3.0.3 Problem 4.5.c

```
[19]: fig, ax = plt.subplots(figsize = (10, 20))  
fig = sm.graphics.plot_partregress_grid(results, fig = fig)  
plt.show()  
print('--> Intercept, x1 have large influence. x2 minimal influence. <--')  
print('--> x3 and above appear to not have any influence <--')
```

Partial Regression Plot



```
--> Intercept, x1 have large influence. x2 minimal influence. <--
--> x3 and above appear to not have any influence <--
```

### 3.0.4 Problem 4.5.d

```
[20]: infl = results.get_influence()
print(infl.summary_table())
print('--> "Obs 15" (i.e. observation 16 because table is 0-based) <--')
print('--> has larger absolute value than others, likely outlier <--')
```

```
=====
=====
```

	obs	endog	fitted	Cook's	student.	hat diag	dffits
ext.stud.		dffits					
			value	d	residual		internal
residual							
	0	29.500	29.429	0.000	0.032	0.441	0.029
0.031		0.028					
	1	27.900	28.061	0.000	-0.063	0.244	-0.036
-0.060		-0.034					
	2	25.900	27.989	0.014	-0.783	0.181	-0.368
-0.771		-0.363					
	3	29.900	29.160	0.001	0.265	0.099	0.088
0.256		0.085					
	4	29.900	26.028	0.142	1.627	0.349	1.190
1.741		1.274					
	5	30.900	32.120	0.018	-0.532	0.395	-0.430
-0.518		-0.418					
	6	28.900	30.362	0.232	-0.943	0.723	-1.525
-0.939		-1.518					
	7	35.900	33.595	0.029	0.909	0.260	0.539
0.903		0.535					
	8	31.500	31.321	0.000	0.076	0.362	0.057
0.073		0.055					
	9	31.000	33.512	0.082	-1.102	0.403	-0.905
-1.111		-0.913					
	10	30.900	33.599	0.057	-1.107	0.316	-0.753
-1.117		-0.760					
	11	30.000	30.165	0.001	-0.082	0.538	-0.089
-0.079		-0.085					
	12	36.900	39.861	0.127	-1.320	0.421	-1.127
-1.360		-1.160					

	13	41.900	40.071	0.283	1.119	0.693	1.681
1.130	1.698						
	14	40.500	39.921	0.005	0.263	0.442	0.234
0.254	0.226						
	15	43.900	39.647	0.372	2.002	0.481	1.929
2.284	2.201						
	16	37.500	35.116	0.156	1.179	0.530	1.250
1.197	1.270						
	17	37.900	39.811	0.062	-0.872	0.448	-0.785
-0.864	-0.778						
	18	44.500	43.609	0.010	0.389	0.398	0.316
0.377	0.307						
	19	37.900	36.340	0.035	0.695	0.420	0.591
0.681	0.580						
	20	38.900	42.620	0.107	-1.525	0.316	-1.036
-1.609	-1.094						
	21	36.900	39.705	0.210	-1.380	0.525	-1.450
-1.431	-1.503						
	22	45.800	42.066	0.097	1.511	0.297	0.983
1.591	1.035						
	23	25.900	26.591	0.049	-0.441	0.717	-0.703
-0.428	-0.682						

```
=====
--> "Obs 15" (i.e. observation 16 because table is 0-based) <--
--> has larger absolute value than others, likely outlier <--
```

## 4 Problem 4.17

```
[21]: df = pd.read_excel('Data/data-table-B10.xlsx')
y, X = patsy.dmatrices('y ~ x1 + x2', df)
model = sm.OLS(y, X)
results = model.fit()
results.model.data.design_info = X.design_info
```


### 4.0.1 Problem 4.17.a

```
[22]: # Get residuals and probability for plot
residuals = results.resid
Prob = [(i - 1/2) / len(y) for i in range(len(y))]

# Calculate OLS using residuals to plot straight line. Get y values from model
resid_results = sm.OLS(Prob, sm.add_constant(sorted(residuals))).fit()
X_range = np.linspace(min(residuals), max(residuals), len(residuals))
```

```
# Normal Probability Plot + straight line
fig, ax = plt.subplots()
ax.scatter(sorted(residuals), Prob)
ax.plot(X_range, resid_results.params[0] + resid_results.params[1] * X_range)
ax.set_xlabel('Residual')
ax.set_ylabel('Probability')
ax.set_ylim(0, 1)
plt.title('Normal Probability Plot')
plt.show()

print('--> Normality plot exhibits negative skew, therefore problems <--')
```

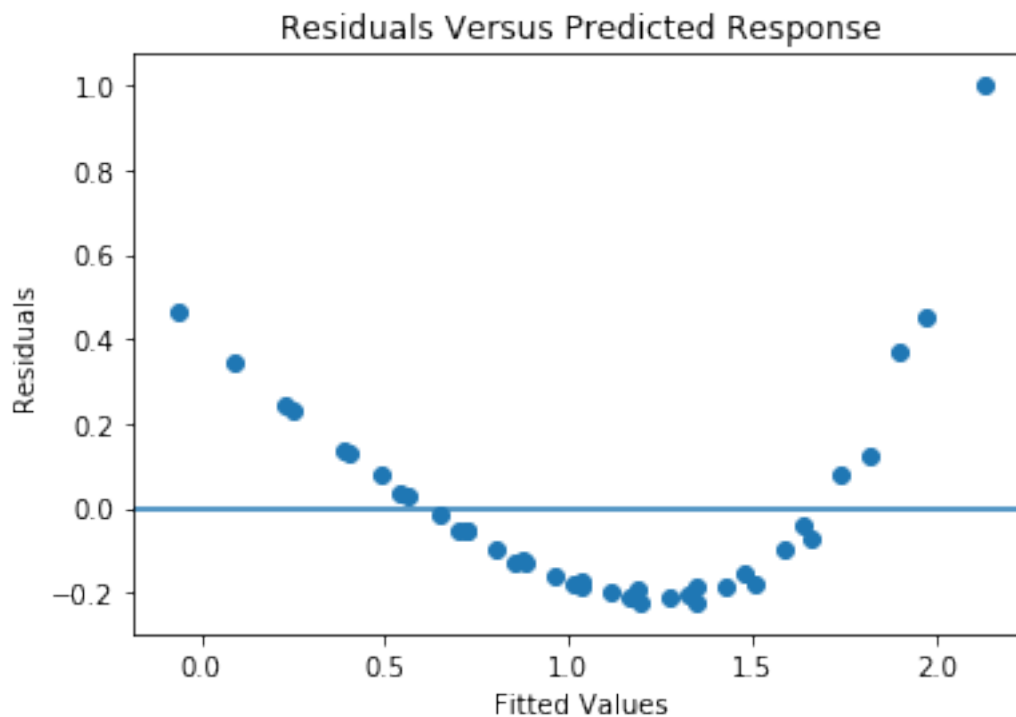


--> Normality plot exhibits negative skew, therefore problems <--

#### 4.0.2 Problem 4.17.b

```
[23]: fig, ax = plt.subplots()
ax.scatter(results.fittedvalues, residuals)
ax.axhline(0)
ax.set_xlabel('Fitted Values')
ax.set_ylabel('Residuals')
plt.title('Residuals Versus Predicted Response')
plt.show()
```

```
print('--> Definite non-linear pattern for residual vs. predicted plot <--')
```



--> Definite non-linear pattern for residual vs. predicted plot <--

#### 4.0.3 Problem 4.17.c

```
[24]: y2, X2 = patsy.dmatrices('y ~ x2', df)
model2 = sm.OLS(y2, X2)
results2 = model2.fit()
results2.model.data.design_info = X2.design_info

infl1 = results.get_influence()
infl2 = results2.get_influence()
press1 = infl1.resid_press
press2 = infl2.resid_press
r_squared_pred1 = 1 - \
    infl1.ess_press/sm.stats.anova_lm(results, typ = 1).sum_sq.sum()
r_squared_pred2 = 1 - \
    infl2.ess_press/sm.stats.anova_lm(results2, typ = 1).sum_sq.sum()

comp_press = pd.DataFrame({'Ordinary Residuals': residuals,
                           'Full Model (PRESS_1)': press1,
```

```
'Viscosity to Temp Only (PRESS_2)': press2})
```

```
[25]: print(comp_press)
```

	Ordinary Residuals	Full Model (PRESS_1)	Viscosity to Temp Only (PRESS_2)
0	0.999076	1.145488	1.506526
1	0.454365	0.506862	0.886794
2	0.123653	0.135198	0.524019
3	-0.074058	-0.079914	0.313249
4	-0.178770	-0.191651	0.203821
5	-0.221481	-0.237440	0.159980
6	-0.221793	-0.239329	0.160660
7	-0.187604	-0.205120	0.198456
8	-0.130516	-0.145596	0.263173
9	-0.055227	-0.063321	0.352562
10	0.372160	0.410135	0.567421
11	0.077448	0.083130	0.238523
12	-0.096263	-0.101347	0.052291
13	-0.182975	-0.190211	-0.037924
14	-0.210686	-0.217648	-0.066133
15	-0.200398	-0.207019	-0.055572
16	-0.159609	-0.165922	-0.013791
17	-0.094721	-0.099723	0.053905
18	-0.013032	-0.013988	0.142049
19	0.080256	0.088446	0.247925
20	-0.042795	-0.047041	-0.173574
21	-0.155507	-0.166500	-0.289264
22	-0.205218	-0.215530	-0.335759
23	-0.209330	-0.217084	-0.335802
24	-0.180441	-0.185957	-0.304061
25	-0.126153	-0.130009	-0.248337
26	-0.051064	-0.052956	-0.172336
27	0.036624	0.038465	-0.082802
28	0.136413	0.146056	0.021991
29	0.244301	0.268542	0.140661
30	-0.182774	-0.211246	-0.646409
31	-0.194985	-0.219217	-0.642721
32	-0.171097	-0.188507	-0.605512
33	-0.122608	-0.133304	-0.547849
34	-0.052320	-0.056512	-0.472294
35	0.031669	0.034206	-0.386085
36	0.127757	0.138902	-0.289256
37	0.230646	0.254115	-0.185306
38	0.342634	0.385215	-0.069492
39	0.464422	0.536768	0.061965



```
[26]: print('Full Model: PRESS Statistic = {} | R^2 (Pred) = {}%'.\
        format(round(infl1.ess_press, 3),
               round(100 * r_squared_pred1, 2)))
print('Partial Model: PRESS Statistic = {} | R^2 (Pred) = {}%'.\
      format(round(infl2.ess_press, 3),
             round(100 * r_squared_pred2, 2)))
print('--> Full Model: Lower PRESS, higher R^2 implies better prediction <--')
```

Full Model: PRESS Statistic = 3.112 | R<sup>2</sup> (Pred) = 77.75%

Partial Model: PRESS Statistic = 6.777 | R<sup>2</sup> (Pred) = 51.54%

--> Full Model: Lower PRESS, higher R<sup>2</sup> implies better prediction <--