

Jacob Miller - Homework 5

March 31, 2020

0.1 Setup

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import patsy
import statsmodels.api as sm
from itertools import combinations, compress
```

```
[2]: def title_print(text):
    '''
    Used throughout to print section titles
    '''
    print()
    print('#' * (len(text) + 4))
    print('|', text, '|')
    print('#' * (len(text) + 4))
```

0.2 Problem 5.1

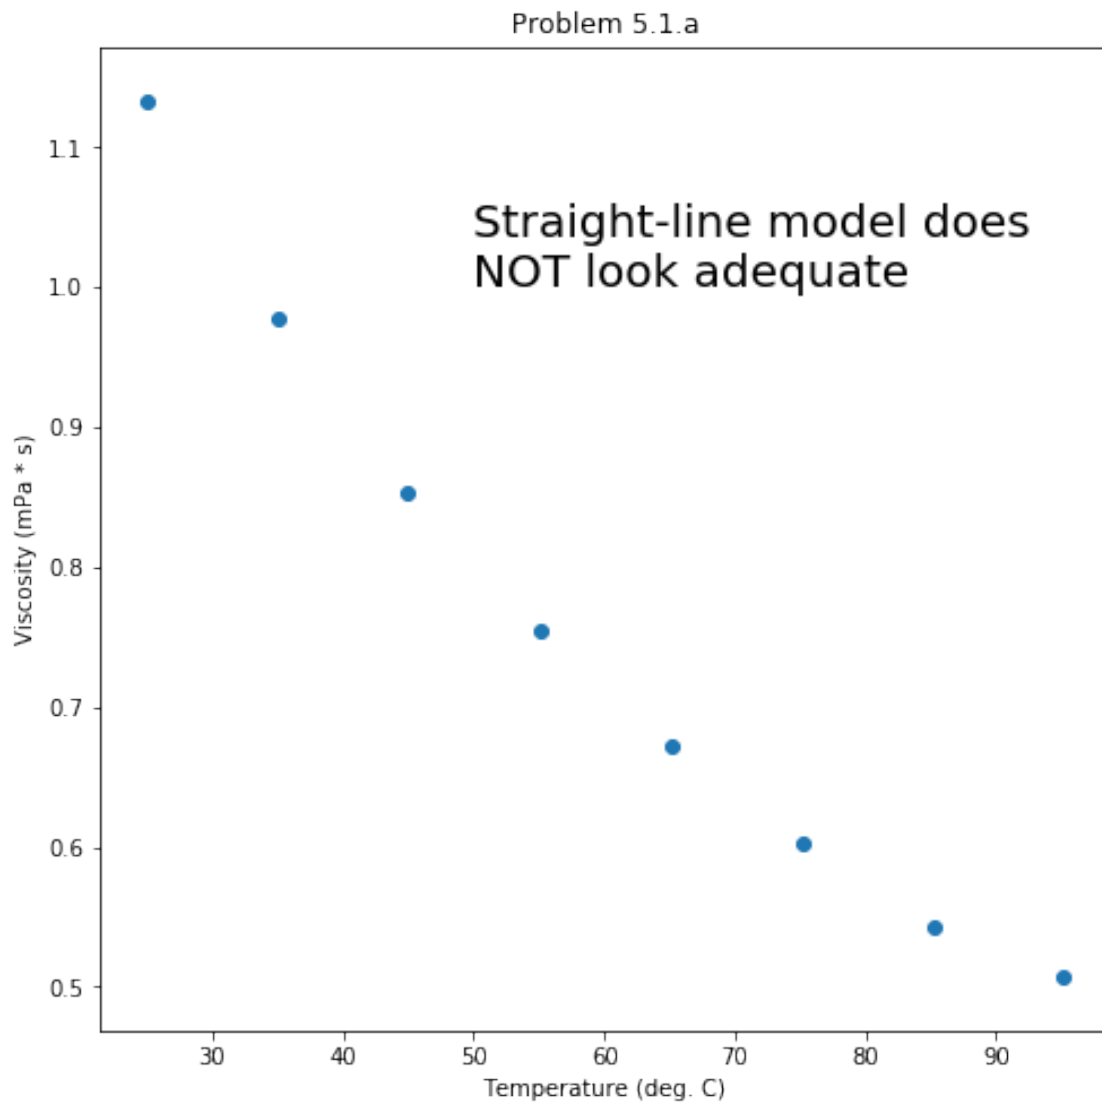
```
[3]: df = pd.DataFrame(data = {'Temperature' : [24.9, 35.0, 44.9, 55.1,
                                                65.2, 75.2, 85.2, 95.2],
                              'Viscosity' : [1.133, 0.9772, 0.8532, 0.7550,
                                              0.6723, 0.6021, 0.5420, 0.5074]})
```

0.2.1 Problem 5.1.a

```
[6]: title_print('Problem 5.1.a')
fig = plt.figure(figsize = (8, 8))
plt.scatter(df['Temperature'], df['Viscosity'])
plt.xlabel('Temperature (deg. C)')
plt.ylabel('Viscosity (mPa * s)')
plt.title('Problem 5.1.a')
plt.text(50, 1, 'Straight-line model does\nNOT look adequate',
        fontdict = {'fontsize': 20})
```

```
plt.show()
```

```
#####  
| Problem 5.1.a |  
#####
```



```
#####  
| Problem 5.1.b |  
#####
```

OLS Regression Results

=====

```

Dep. Variable:          Viscosity    R-squared:          0.960
Model:                  OLS          Adj. R-squared:      0.954
Method:                 Least Squares  F-statistic:        144.6
Date:                   Tue, 31 Mar 2020  Prob (F-statistic):  2.01e-05
Time:                   15:02:37      Log-Likelihood:     14.187
No. Observations:      8             AIC:                -24.37
Df Residuals:          6             BIC:                -24.21
Df Model:               1
Covariance Type:       nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      1.2815      0.047      27.343      0.000       1.167       1.396
Temperature    -0.0088      0.001     -12.024      0.000      -0.011      -0.007
=====
Omnibus:                1.431    Durbin-Watson:           0.734
Prob(Omnibus):          0.489    Jarque-Bera (JB):       0.942
Skew:                   0.642    Prob(JB):               0.624
Kurtosis:               1.915    Cond. No.               180.
=====

```

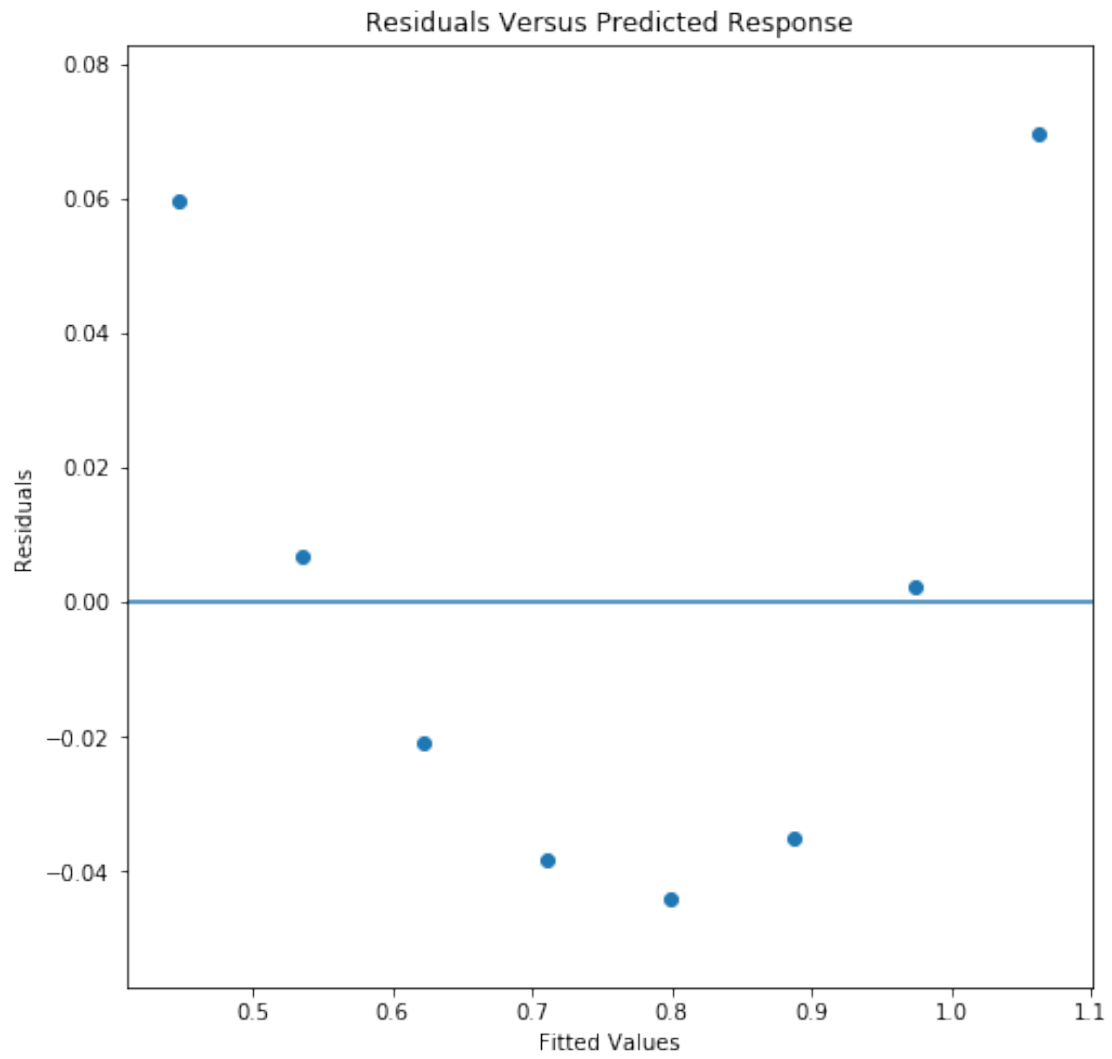
Warnings:


```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
```

```

/Users/Jake/anaconda3/envs/Data/lib/python3.7/site-
packages/scipy/stats/stats.py:1535: UserWarning: kurtosistest only valid for
n>=20 ... continuing anyway, n=8
  "anyway, n=%i" % int(n))

```





```
--> Clear non-linearity in residual plot <--
--> Normality appears to have problems <--
```

```
#####
| Problem 5.1.c |
#####
```

OLS Regression Results

```
=====
Dep. Variable:          Viscosity    R-squared:                0.999
Model:                  OLS          Adj. R-squared:           0.999
Method:                 Least Squares  F-statistic:              7952.
Date:                   Tue, 31 Mar 2020  Prob (F-statistic):      1.34e-10
```

```

Time:                  15:02:37    Log-Likelihood:          30.056
No. Observations:      8    AIC:                -56.11
Df Residuals:          6    BIC:                -55.95
Df Model:              1
Covariance Type:      nonrobust

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept          2.6651      0.022    123.721      0.000      2.612
2.718
np.log(Temperature) -0.4762      0.005    -89.172      0.000     -0.489
-0.463
=====
Omnibus:            3.567    Durbin-Watson:           1.623
Prob(Omnibus):      0.168    Jarque-Bera (JB):           1.267
Skew:               0.975    Prob(JB):                 0.531
Kurtosis:           2.954    Cond. No.                  40.0
=====

```


Warnings:

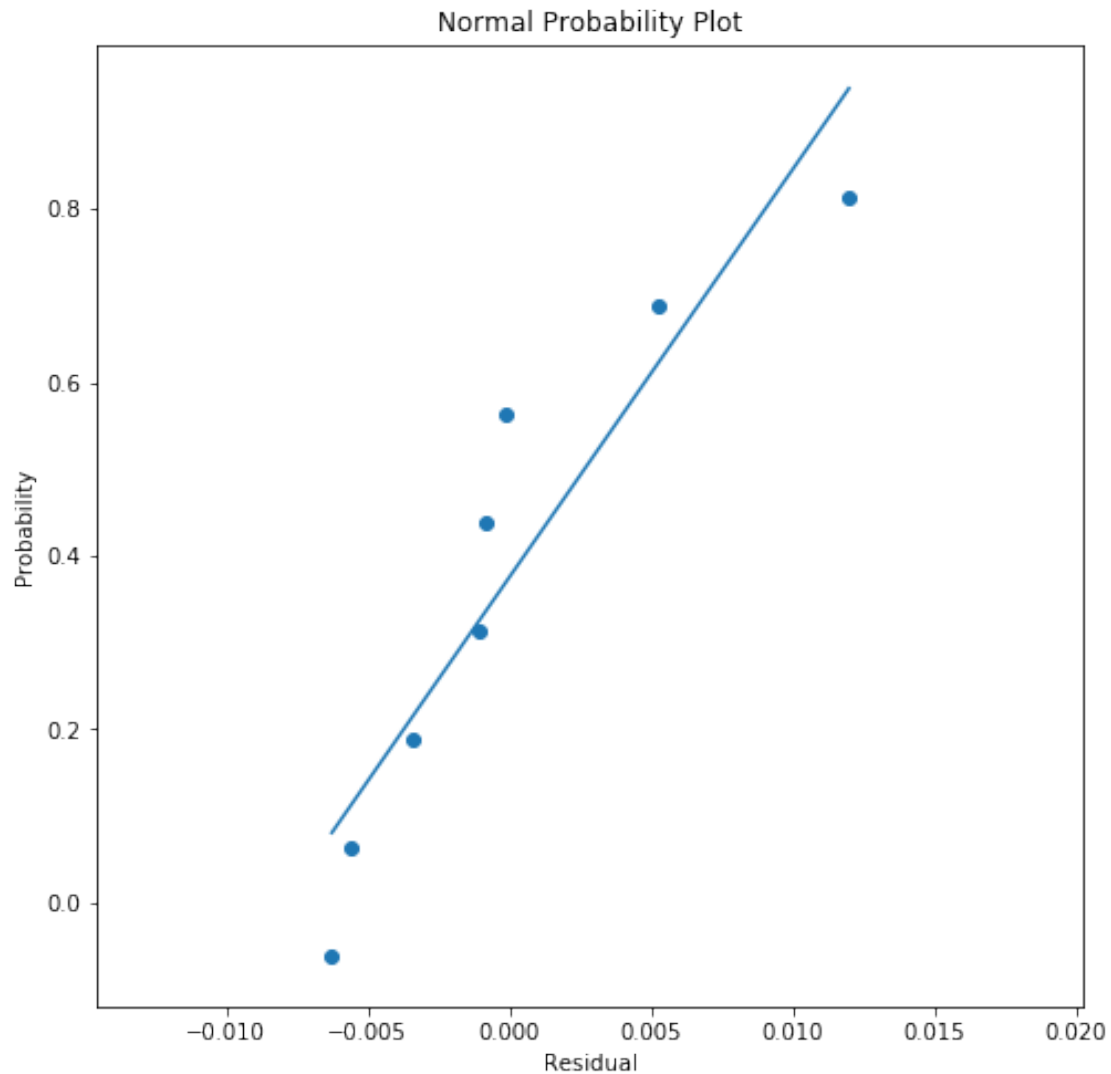
```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
```

```

/Users/Jake/anaconda3/envs/Data/lib/python3.7/site-
packages/scipy/stats/stats.py:1535: UserWarning: kurtosistest only valid for
n>=20 ... continuing anyway, n=8
"anyway, n=%i" % int(n))

```





```
--> Residual vs. Response plot mildly improved <--  
--> R-squared value slightly increased <--  
--> Overall improvement seems minimal <--
```

0.2.2 Problem 5.1.b

```
[7]: title_print('Problem 5.1.b')  
y, X = patsy.dmatrices('Viscosity ~ Temperature', df)  
model = sm.OLS(y, X)  
results = model.fit()  
results.model.data.design_info = X.design_info  
print(results.summary())
```



```

# Get residuals and probability for plot
residuals = results.resid
Prob = [(i - 1/2) / len(y) for i in range(len(y))]

# Plot residuals vs. fitted values
fig, ax = plt.subplots(figsize = (8, 8))
ax.scatter(results.fittedvalues, residuals)
ax.axhline(0)
ax.set_xlabel('Fitted Values')
ax.set_ylabel('Residuals')
plt.title('Residuals Versus Predicted Response')
plt.show()

# Calculate OLS using residuals to plot straight line. Get y values from model
resid_results = sm.OLS(Prob, sm.add_constant(sorted(residuals))).fit()
X_range = np.linspace(min(residuals), max(residuals), len(residuals))

# Normality plot
fig = plt.figure(figsize = (8, 8))
plt.scatter(sorted(residuals), Prob)
plt.plot(X_range, resid_results.params[0] + resid_results.params[1] * X_range)
plt.xlabel('Residual')
plt.ylabel('Probability')
plt.title('Normal Probability Plot')
plt.show()

print('\n--> Clear non-linearity in residual plot <--')
print('--> Normality appears to have problems <--')

```

```

#####
| Problem 5.1.b |
#####

```

OLS Regression Results					
=====					
Dep. Variable:	Viscosity	R-squared:	0.960		
Model:	OLS	Adj. R-squared:	0.954		
Method:	Least Squares	F-statistic:	144.6		
Date:	Tue, 31 Mar 2020	Prob (F-statistic):	2.01e-05		
Time:	15:03:43	Log-Likelihood:	14.187		
No. Observations:	8	AIC:	-24.37		
Df Residuals:	6	BIC:	-24.21		
Df Model:	1				
Covariance Type:	nonrobust				
=====					
	coef	std err	t	P> t	[0.025 0.975]

```

-----
Intercept      1.2815      0.047      27.343      0.000      1.167      1.396
Temperature    -0.0088      0.001     -12.024      0.000     -0.011     -0.007
=====
Omnibus:                1.431   Durbin-Watson:                0.734
Prob(Omnibus):          0.489   Jarque-Bera (JB):          0.942
Skew:                   0.642   Prob(JB):                  0.624
Kurtosis:               1.915   Cond. No.                  180.
=====

```


Warnings:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

/Users/Jake/anaconda3/envs/Data/lib/python3.7/site-
packages/scipy/stats/stats.py:1535: UserWarning: kurtosistest only valid for
n>=20 ... continuing anyway, n=8
  "anyway, n=%i" % int(n))

```





--> Clear non-linearity in residual plot <--
--> Normality appears to have problems <--

0.2.3 Problem 5.1.c

```
[ ]: title_print('Problem 5.1.c')
y, X = patsy.dmatrices('Viscosity ~ np.log(Temperature)', df)
model = sm.OLS(y, X)
results = model.fit()
results.model.data.design_info = X.design_info
print(results.summary())
```

```

# Get residuals and probability for plot
residuals = results.resid
Prob = [(i - 1/2) / len(y) for i in range(len(y))]

# Plot residuals vs. fitted values
fig, ax = plt.subplots(figsize = (8, 8))
ax.scatter(results.fittedvalues, residuals)
ax.axhline(0)
ax.set_xlabel('Fitted Values')
ax.set_ylabel('Residuals')
plt.title('Residuals Versus Predicted Response')
plt.show()

# Calculate OLS using residuals to plot straight line. Get y values from model
resid_results = sm.OLS(Prob, sm.add_constant(sorted(residuals))).fit()
X_range = np.linspace(min(residuals), max(residuals), len(residuals))

# Normality plot
fig = plt.figure(figsize = (8, 8))
plt.scatter(sorted(residuals), Prob)
plt.plot(X_range, resid_results.params[0] + resid_results.params[1] * X_range)
plt.xlabel('Residual')
plt.ylabel('Probability')
plt.title('Normal Probability Plot')
plt.show()

print('\n--> Residual vs. Response plot mildly improved <--')
print('--> R-squared value slightly increased <--')
print('--> Overall improvement seems minimal <--')

```

0.3 Problem 6.1

Note that Python uses 0-based indexing, so add 1 to the points in the following dataframe

```

[14]: def run_analysis(drop_point = None):
    """
    Parameters
    -----
    drop_point : int or list-like, optional
        Point(s) to drop for analysis. The default is None.

    Returns
    -----
    Points to potentially drop, if drop_point == None.
    """

```

```

'''
df = pd.read_excel('Data/data-table-b2.xlsx')
df.columns = ['y', 'x1', 'x2', 'x3', 'x4', 'x5']

# Drop influential points, if necessary
if drop_point:
    df = df.drop(np.array(drop_point) - 1)
y, X = patsy.dmatrices('y ~ x1 + x2 + x3 + x4 + x5', df)

# Fit model, get influence statistics
model = sm.OLS(y, X)
results = model.fit()
results.model.data.design_info = X.design_info

# If dropping points, only run to here and then exit function
if drop_point:
    print('\nPoints dropped: {}'.format(drop_point))
    print('Coefficients: {}'.format(np.round(results.params, 3)))
    print('R-squared: {}'.format(round(results.rsquared, 3)))
    return
else:
    print('Coefficients: {}'.format(np.round(results.params, 3)))
    print('R-squared: {}'.format(round(results.rsquared, 3)))

infl = results.get_influence()
infl_df = infl.summary_frame()
print(infl_df)

# Hat diagonals
n, p = df.shape
lev_pt = 2 * p / n
dhat_pts = list(infl_df[infl_df['hat_diag'] > lev_pt].index + 1)
print('\n***| Hat Diagonal |***')
print('Leverage calculation (2 * p \ n) = {}'.format(round(lev_pt, 3)))
print('Points where hat diagonal exceeds leverage calculation: {}'.format(dhat_pts))

# Cook's D
print('\n***| Cook\'s D |***')
print('Points where Cook\'s D is > 1: {}'.format(list(infl_df[infl_df['cooks_d'] > 1].index + 1)))

# DFFITS
print('\n***| DFFITS |***')
DFFITS_cutoff = 2 * np.sqrt(p / n)
print('Points which exceed DFFITS cutoff: {}'.format(list(infl_df[infl_df['dffits'] > DFFITS_cutoff].index + 1)))

```

```

# DFBETAS
print('\n***| DFBETAS |***')
DFBETAS_cutoff = 2 / np.sqrt(n)
for col in infl_df.columns:
    if 'dfb' in col:
        print('Points which exceed DFBETAS cutoff for {}: {}'.format(col,
            list(infl_df[infl_df[col] > DFBETAS_cutoff].index + 1)))

# COVRATIO
print('\n***| COVRATIO |***')
COVRATIO_cutoff_pos = 1 + 3 * p / n
COVRATIO_cutoff_neg = 1 - 3 * p / n
gt_cutoff = np.array(list(compress(range(len(infl.cov_ratio)),
    infl.cov_ratio > COVRATIO_cutoff_pos))) + 1
lt_cutoff = np.array(list(compress(range(len(infl.cov_ratio)),
    infl.cov_ratio < COVRATIO_cutoff_neg))) + 1
print('Points which are greater than COVRATIO upper bound cutoff: {}'.format(gt_cutoff))
print('Points which are less than COVRATIO lower bound cutoff: {}'.format(lt_cutoff))

return dhat_pts

```

```

[15]: # All points
leverage_points = run_analysis(drop_point = None)

title_print('Try dropping influential points')

# Drop influential points
for i in range(1, len(leverage_points) + 1):
    comb = combinations(leverage_points, i)
    for pts in comb:
        run_analysis(pts)

# Points 1, 4 appear influential
print('\n--> Points 1 and 4 appear influential <--')

```

Coefficients: [3.25436e+02 6.80000e-02 2.55200e+00 3.80000e+00 -2.29490e+01
2.41700e+00]

R-squared: 0.899

	dfb_Intercept	dfb_x1	dfb_x2	dfb_x3	dfb_x4	dfb_x5	cooks_d	\
0	-0.650262	-0.435297	0.193029	1.411300	0.378643	-1.099424	0.375089	
1	-0.011559	0.005656	0.024343	0.003432	-0.010825	0.011227	0.000309	

2	-0.014808	-0.024453	-0.016779	0.023476	0.022013	0.026848	0.005696
3	0.002208	0.987087	-1.022694	-0.878006	1.873989	-0.615708	1.101218
4	0.040237	-0.059983	-0.032735	0.011278	-0.045723	0.012378	0.001509
5	-0.008331	0.077484	-0.012171	-0.036789	0.029652	0.012072	0.002200
6	0.034348	-0.251545	0.004691	0.107034	-0.025644	-0.107824	0.015202
7	0.314302	-0.758889	-0.255013	0.104827	-0.034604	-0.286666	0.129157
8	0.181243	0.176124	-0.192785	-0.209058	-0.326112	0.513198	0.077888
9	0.067652	0.078741	-0.083447	-0.093096	-0.127054	0.239105	0.021118
10	-0.002020	-0.001250	0.002395	0.004031	0.001818	-0.006794	0.000025
11	-0.043673	0.020632	0.010777	-0.048159	0.092226	0.053902	0.014138
12	-0.016456	-0.011500	0.027630	-0.012353	0.029844	-0.006728	0.001216
13	0.022355	-0.004016	-0.034586	0.003643	-0.018740	0.005802	0.000404
14	0.023058	0.001983	0.003856	-0.065121	0.010484	0.015792	0.002367
15	0.438398	-0.138762	-0.625543	-0.143128	-0.196630	0.205495	0.085585
16	-0.021730	0.010100	0.015312	0.004751	0.020566	0.001969	0.000234
17	-0.058175	0.082987	0.292622	0.182041	-0.312095	-0.210324	0.153429
18	0.346051	-0.017083	-0.196286	-0.273684	-0.382732	0.311561	0.034934
19	-0.086891	0.010615	0.014062	0.066022	0.144261	-0.106660	0.005143
20	-0.011577	0.058752	0.076128	-0.052995	-0.076632	0.119400	0.006715
21	-0.942814	0.434040	1.116032	0.144539	0.496061	0.128682	0.270561
22	-0.010436	0.026076	-0.007384	-0.061900	0.047793	0.081672	0.006697
23	0.049658	0.265351	0.044671	0.220436	-0.480786	-0.121384	0.245601
24	-0.023504	0.034986	0.011158	0.006817	0.004580	0.017579	0.000459
25	0.126933	-0.156130	-0.038363	-0.034092	-0.090438	-0.045799	0.009698
26	0.083988	-0.127925	0.025901	-0.015913	-0.110619	-0.018061	0.008388
27	-0.019469	0.026777	-0.008167	0.000589	0.036939	-0.004132	0.000549
28	-0.206022	-0.122063	0.342483	0.205125	-0.010703	-0.098899	0.036763

	standard_resid	hat_diag	dffits_internal	student_resid	dffits
0	0.969344	0.705461	1.500179	0.968017	1.498124
1	0.129202	0.099829	0.043027	0.126408	0.042096
2	0.561851	0.097692	0.184873	0.553311	0.182063
3	2.894219	0.440963	2.570468	3.549908	3.152811
4	-0.205319	0.176817	-0.095157	-0.200990	-0.093151
5	0.266154	0.157045	0.114880	0.260706	0.112528
6	-0.567502	0.220710	-0.302015	-0.558955	-0.297466
7	-1.636853	0.224346	-0.880308	-1.703144	-0.915960
8	-1.515296	0.169111	-0.683615	-1.562007	-0.704689
9	-0.953275	0.122370	-0.355958	-0.951303	-0.355222
10	0.043494	0.074643	0.012353	0.042540	0.012082
11	-1.114583	0.063917	-0.291249	-1.120772	-0.292866
12	-0.272683	0.089362	-0.085420	-0.267121	-0.083678
13	0.089136	0.233600	0.049211	0.087192	0.048138
14	-0.350873	0.103433	-0.119176	-0.344083	-0.116870
15	0.857663	0.411107	0.716598	0.852554	0.712330
16	0.085527	0.161091	0.037478	0.083660	0.036661
17	-1.905844	0.202199	-0.959466	-2.031229	-1.022590
18	0.799950	0.246732	0.457826	0.793483	0.454125

19	-0.379522	0.176443	-0.175668	-0.372347	-0.172347
20	0.518279	0.130428	0.200722	0.509873	0.197467
21	2.091182	0.270723	1.274115	2.272648	1.384678
22	0.480335	0.148327	0.200455	0.472151	0.197040
23	-1.799757	0.312686	-1.213921	-1.898987	-1.280851
24	-0.071693	0.348858	-0.052476	-0.070125	-0.051329
25	0.521723	0.176124	0.241223	0.513302	0.237329
26	0.529176	0.152345	0.224339	0.520724	0.220756
27	-0.151075	0.126199	-0.057414	-0.147828	-0.056179
28	1.086494	0.157438	0.469658	1.090978	0.471596

| Hat Diagonal |

Leverage calculation ($2 * p \setminus n$) = 0.414

Points where hat diagonal exceeds leverage calculation: [1, 4]

| Cook's D |

Points where Cook's D is > 1: [4]

| DFFITS |

Points which exceed DFFITS cutoff: [1, 4, 22]

| DFBETAS |

Points which exceed DFBETAS cutoff for dfb_Intercept: [16]

Points which exceed DFBETAS cutoff for dfb_x1: [4, 22]

Points which exceed DFBETAS cutoff for dfb_x2: [22]

Points which exceed DFBETAS cutoff for dfb_x3: [1]

Points which exceed DFBETAS cutoff for dfb_x4: [1, 4, 22]

Points which exceed DFBETAS cutoff for dfb_x5: [9]

| COVRATIO |

Points which are greater than COVRATIO upper bound cutoff: [1 14 16 25]

Points which are less than COVRATIO lower bound cutoff: [4]

#####

| Try dropping influential points |

#####

Points dropped: (1,)

Coefficients: [3.8803e+02 8.0000e-02 2.3110e+00 1.7350e+00 -2.3975e+01
4.4080e+00]

R-squared: 0.899

Points dropped: (4,)

Coefficients: [3.25263e+02 4.40000e-02 3.59300e+00 4.84600e+00 -2.70800e+01
3.32500e+00]

R-squared: 0.934

Points dropped: (1, 4)

```
Coefficients: [ 3.58769e+02  5.20000e-02  3.43400e+00  3.71100e+00 -2.75100e+01  
               4.36500e+00]  
R-squared: 0.933
```

```
--> Points 1 and 4 appear influential <--
```